

Implementation Details for Algorithms:

- I used the recommended fold value of $k = 10$
- I used the max depth of the algorithm to 10.
- For the min size before split ie. remaining n values after I don't split, is set to 5.
- Minimal Gain is set to 0.01 being close to 0 but not 0.
- The bootstrap ratio has been set to 0.1 so 10% of the data got resampled.
- Stopping criterion like minimal size for split criterion, maximal depth, minimal_gain

For the Wine Dataset using ID3:

1 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.951

Precision: 0.939

Recall: 0.929

F-Score (beta=1): 0.929

5 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.974

Precision: 0.968

Recall: 0.962

F-Score(beta=1): 0.962

10 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.982

Precision: 0.974

Recall: 0.976

F-Score(beta=1): 0.973

20 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.985

Precision: 0.98

Recall: 0.979

F-Score(beta=1): 0.979

30 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.981

Precision: 0.975

Recall: 0.973

F-Score(beta=1): 0.971

40 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.989

Precision: 0.985

Recall: 0.984

F-Score(beta=1): 0.984


50 Trees Random Forest of Wine Dataset with ID3:

Accuracy: 0.985

Precision: 0.979

Recall: 0.98

F-Score(beta=1): 0.978



For the House Dataset using ID3:

1 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.915

Precision: 0.914

Recall: 0.91

F-Score(beta=1): 0.909

5 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.94

Precision: 0.94

Recall: 0.937

F-Score(beta=1): 0.937

10 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.956

Precision: 0.957

Recall: 0.953

F-Score(beta=1): 0.954

20 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.954

Precision: 0.951

Recall: 0.954

F-Score(beta=1): 0.952

30 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.961

Precision: 0.958

Recall: 0.961

F-Score(beta=1): 0.959

40 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.956

Precision: 0.953

Recall: 0.957

F-Score(beta=1): 0.954


50 Trees Random Forest of House Dataset with ID3:

Accuracy: 0.961

Precision: 0.959

Recall: 0.961

F-Score(beta=1): 0.959



For the Wine Dataset using Gini:

1 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.938

Precision: 0.922

Recall: 0.908

F-Score(beta=1): 0.909

5 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.974

Precision: 0.968

Recall: 0.958

F-Score(beta=1): 0.96

10 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.982

Precision: 0.976

Recall: 0.975

F-Score(beta=1): 0.974

20 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.985

Precision: 0.981

Recall: 0.979

F-Score(beta=1): 0.979

30 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.985

Precision: 0.979

Recall: 0.98

F-Score(beta=1): 0.978

40 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.984

Precision: 0.98

Recall: 0.981

F-Score(beta=1): 0.978

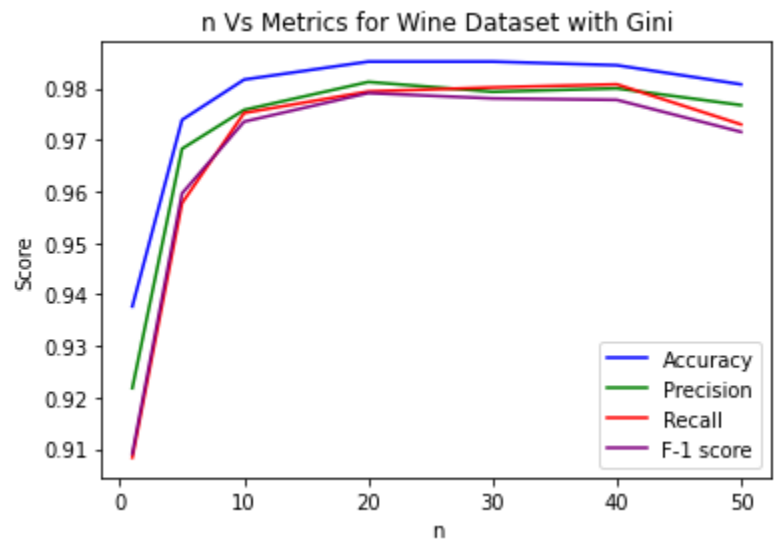
50 Trees Random Forest of Wine Dataset with Gini:

Accuracy: 0.981

Precision: 0.977

Recall: 0.973

F-Score(beta=1): 0.972



For the House Dataset using Gini:

1 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.931

Precision: 0.927

Recall: 0.93

F-Score(beta=1): 0.926

5 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.945

Precision: 0.942

Recall: 0.944

F-Score(beta=1): 0.942

10 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.947

Precision: 0.949

Recall: 0.944

F-Score(beta=1): 0.944

20 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.954

Precision: 0.952

Recall: 0.955

F-Score(beta=1): 0.952

30 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.956

Precision: 0.953

Recall: 0.957

F-Score(beta=1): 0.954

40 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.959

Precision: 0.959

Recall: 0.957

F-Score(beta=1): 0.956


50 Trees Random Forest of House Dataset with Gini:

Accuracy: 0.959

Precision: 0.957

Recall: 0.959

F-Score(beta=1): 0.957



4) For each metric being evaluated (and for each dataset), discuss which value of *ntree* you would select if you were to deploy this classifier in real life. Explain your reasoning.

a) For the Wine Dataset:

- i) Accuracy: We observe that while using Entropy, the value *ntree* = 20 returns the most optimal results. In the case of Cart, the value *ntree* = 40 returns the most optimal results. Hence, backed by our findings, I would choose the value of *ntree* to be 40 if I were to deploy this classifier in real life for the most optimal accuracy results.
- ii) Precision: We observe that while using Entropy, the value *ntree* = 40 returns the most optimal results. In case of Cart, the range of *ntree* value from 20 - 40, almost remains constant showing a next to none loss in precision value with a standard deviation of ± 0.001 . Hence, backed by our findings, I would choose the value of *ntree* to be 40 if I were to deploy this classifier in real life for the most optimal precision results.
- iii) Recall: We observe that while using Entropy, the value *ntree* = 40 returns the most optimal results. In the case of Cart, the range of *ntree* value from 20 - 40, almost remains constant showing a next to none loss in recall value with a standard deviation of ± 0.001 . Hence, backed by our findings, I would choose the value of *ntree* to be 40 if I were to deploy this classifier in real life for the most optimal recall results.
- iv) F-1: We observe that while using Entropy, the value *ntree* = 40 returns the most optimal results. In the case of Cart, the range of *ntree* value from 20 - 40, almost remains constant showing a next to none loss in F-1 value. Hence, backed by our findings, I would choose the value of *ntree* to be 40 if I were to deploy this classifier in real life for the most optimal F-1 results.

Hence, overall I would pick *ntree* value to be 40 for the most optimal results across all metrics if I were to deploy the model in real life.

b) For the House Votes Dataset:

- i) Accuracy: We observe that while using Entropy, the value *ntree* = 30/50 returns the most optimal results. In the case of Cart, the value *ntree* = 50 returns the most optimal results.
- ii) Precision: We observe that while using Entropy, the value *ntree* = 30/50 returns the most optimal results. In the case of Cart, the value *ntree* = 40 or 50 showing a very little deviation and returns the most optimal results.
- iii) Recall: We observe that while using Entropy, the value *ntree* = 30/50 returns the most optimal results. In the case of Cart, the value *ntree* = 50 returns the most optimal results.
- iv) F-1: We observe that while using Entropy, the value *ntree* = 30/50 returns the most optimal results. In the case of Cart, the value *ntree* = 50 returns the most optimal results.

Hence, overall I would pick *ntree* value to be either 30 or 50 in case of Entropy, and would pick *ntree* value as 50 in case of Cart for the most optimal results across all metrics. Hence, if I were to deploy the model in real life, I would choose *ntree* value as 50 to obtain the most optimal results.

5) *Discuss (on a high level) which metrics were more directly affected by changing the value of ntree and, more generally, how such changes affected the performance of your algorithm. For instance: was the accuracy of the random forest particularly sensitive to increasing ntree past a given value? Was the F1 score a “harder” metric to optimize, possibly requiring a significant number of trees in the ensemble? Is there a point beyond which adding more trees does not improve performance—or makes the performance worse?*

In case of Entropy(ID3), all the metrics were equally sensitive to the value of ntree. This makes sense as accuracy and f1-score are somewhat related to the values of precision and recall. In general, as the value of ntree increased, the value of all the metrics also increased showing an upward trend and slowly stabilizing, as should be the case. Hence the performance of the classifier increased as the value ntree increased.

In the case of the Wine dataset, we observe all the values increasing until ntree = 40 after which they start to dip little as ntree value reaches 50 which could be due to overfitting of the model as the number of trees increases..

However, in the case of House Votes dataset, we observe the values peaking early around ntree = 30 and stabilizing after all the way to ntree = 50 indicating that increasing the number of trees beyond that point may not provide significant improvements in performance.

In the case of Cart(Gini), all the metrics were equally sensitive to the value of ntree. This makes sense as accuracy and f1-score are somewhat related to the values of precision and recall. In general, as the value of ntree increased, the value of all the metrics also increased showing an upward trend and slowly stabilizing, as should be the case. Hence the performance of the classifier increased as the value ntree increased.

In the case of the Wine dataset, we observe all the values increasing until ntree = 40 after which they start to dip little as ntree value reaches 50. However, in the case of House Votes dataset, we observe the values constantly increasing until the value of ntree reaches 50 where we, in our test case, received the most optimal results. We know that As the number of trees increases, the variance of the model decreases, which leads to better generalization performance.

The difference in the behavior of the Wine and House Votes datasets could be explained due to the complexity of the datasets and the size of the dataset. The Wine dataset is less complex and smaller than the House Votes dataset, which is why the performance of the model starts to dip after ntree = 40.

(Extra Points #2: 8 Points) Analyze a third dataset: the Breast Cancer Dataset. The goal, here, is to classify whether tissue removed via a biopsy indicates whether a person may or may not have breast cancer. There are 699 instances in this dataset. Each instance is described by 9 numerical 4 attributes, and there are 2 classes. You should present the same analyses and graphs as discussed above. This dataset can be found in the same zip file as the two main datasets.

For the Cancer Dataset using ID3:

1 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.941

Precision: 0.933

Recall: 0.896

F-Score(beta=1): 0.912

5 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.947

Precision: 0.946

Recall: 0.9

F-Score(beta=1): 0.921

10 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.95

Precision: 0.941

Recall: 0.912

F-Score(beta=1): 0.925

20 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.959

Precision: 0.943

Recall: 0.938

F-Score(beta=1): 0.94

30 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.956

Precision: 0.941

Recall: 0.929

F-Score(beta=1): 0.935

40 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.954

Precision: 0.946

Recall: 0.921


F-Score(beta=1): 0.932

50 Trees Random Forest of Cancer Dataset with ID3:

Accuracy: 0.951

Precision: 0.94

Recall: 0.921



F-Score(beta=1): 0.928

For the Cancer Dataset using Gini:

1 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.923

Precision: 0.923

Recall: 0.851

F-Score(beta=1): 0.881

5 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.95

Precision: 0.953

Recall: 0.901

F-Score(beta=1): 0.925

10 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.958

Precision: 0.946

Recall: 0.934

F-Score(beta=1): 0.94

20 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.949

Precision: 0.938

Recall: 0.913

F-Score(beta=1): 0.924

30 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.961

Precision: 0.954

Recall: 0.938

F-Score(beta=1): 0.944

40 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.959

Precision: 0.946

Recall: 0.934

F-Score(beta=1): 0.939

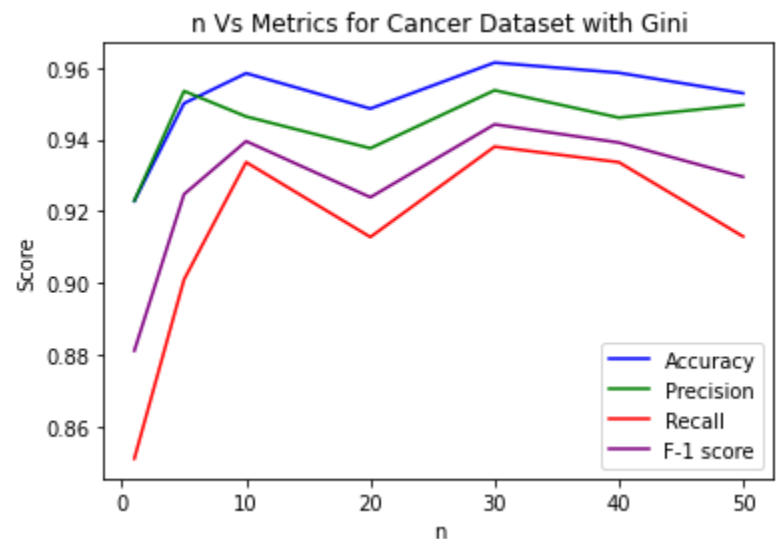
50 Trees Random Forest of Cancer Dataset with Gini:

Accuracy: 0.953

Precision: 0.95

Recall: 0.913

F-Score(beta=1): 0.93



4)

For the Cancer Dataset:

Accuracy: We observe that while using Entropy, the value $n_{tree} = 20$ returns the most optimal results. In the case of Cart, the value $n_{tree} = 30$ returns the most optimal results.

Precision: We observe that while using Entropy, the value $n_{tree} = 5/40$ (with $n_{tree} = 20$ showing a very tiny variance of 0.003) returns the most optimal results. In the case of Cart, the value $n_{tree} = 5/30$ showing a very little deviation and returns the most optimal results.

Recall: We observe that while using Entropy, the value $n_{tree} = 20$ returns the most optimal results. In the case of Cart, the value $n_{tree} = 30$ returns the most optimal results.

F-1: We observe that while using Entropy, the value $n_{tree} = 20$ returns the most optimal results. In the case of Cart, the value $n_{tree} = 30$ returns the most optimal results.

Hence, overall I would pick n_{tree} value to be 20 in case of Entropy, and would pick n_{tree} value as 30 in case of Cart for the most optimal results across all metrics. Hence, if I were to deploy the model in real life, I would choose n_{tree} value between 20 to 30 to obtain the most optimal results.

5)

In both cases ie. Entropy and Cart, we observe that the claes increase rapidly to n_{tree} value of 20/30 and then start diverging both ways but in sort of a stabilizing manner until starting to decrease after reaching $n_{tree} = 50$. I believe this trend is due to the complexity of the dataset and increasing the value of n_{tree} leads to overfitting after a certain point and doesn't lead to much informational gain (around 20-30).

File Descriptions:

- 1) Hw3.py: Runs all the code
- 2) decisionTree.py: implementation of decision tree forest
- 3) randomForest.py. Implementation of random forest
- 4) evaluation.py: Metrics and results